

```
In [1]: from __future__ import absolute_import
        from __future__ import division
        from __future__ import print_function

        import numpy as np
        import scipy.signal as sg
        from utils import *
```

```

In [2]: # superclass of modules
class Module:
    """
    Module is a super class. It could be a single layer, or a multilayer perceptron.
    """

    def __init__(self):
        self.train = True
        return

    def forward(self, _input):
        """
         $h = f(z)$ ;  $z$  is the input, and  $h$  is the output.

        Inputs:
        _input:  $z$ 

        Returns:
        output  $h$ 
        """
        pass

    def backward(self, _input, _gradOutput):
        """
        Compute:
        gradient w.r.t. _input
        gradient w.r.t. trainable parameters

        Inputs:
        _input:  $z$ 
        _gradOutput:  $dL/dh$ 

        Returns:
        gradInput:  $dL/dz$ 
        """
        pass

    def parameters(self):
        """
        Return the value of trainable parameters and its corresponding gradient (Used for gradient descent)

        Returns:
        params, gradParams
        """

```

pass

```
def training(self):  
    """  
    Turn the module into training mode.(Only useful for Dropout Layer)  
    Ignore it if you are not using Dropout.  
    """  
    self.train = True  
  
def evaluate(self):  
    """  
    Turn the module into evaluate mode.(Only useful for Dropout Layer)  
    Ignore it if you are not using Dropout.  
    """  
    self.train = False
```

```

In [3]: class Sequential(Module):
        """
        Sequential provides a way to plug layers together in a feed-forward manner.
        """
        def __init__(self):
            Module.__init__(self)
            self.layers = [] # Layers contain all the layers in order

        def add(self, layer):
            self.layers.append(layer) # Add another layer at the end

        def size(self):
            return len(self.layers) # How many layers.

        def forward(self, _input):
            """
            Feed forward through all the layers, and return the output of the last
            layer
            """
            # self._inputs saves the input of each layer
            # self._inputs[i] is the input of i-th layer
            self._inputs = [_input]

            for i in range(self.size()):
                self._inputs.append(self.layers[i].forward(self._inputs[i]))

            # The last element of self._inputs is the output of last layer
            self._output = self._inputs[-1]
            return self._output

        def backward(self, _input, _gradOutput):
            """
            Backpropagate through all the layers using chain rule.
            """
            # self._gradInputs[i] is the gradient of loss w.r.t. the input of i-th
            layer
            self._gradInputs = [None] * (self.size() + 1)
            self._gradInputs[self.size()] = _gradOutput

```

```

        for i in range(self.size(), 0, -1):
            self._gradInputs[i-1] = self.layers[i-1].backward(self._inputs[i-1], self._gradInputs[i])

        self._gradInput = self._gradInputs[0]
        return self._gradInput

    def parameters(self):
        """
        Return trainable parameters and its corresponding gradient in a nested
list
        """
        params = []
        gradParams = []
        for m in self.layers:
            _p, _g = m.parameters()
            if _p is not None:
                params.append(_p)
                gradParams.append(_g)
        return params, gradParams

    def training(self):
        """
        Turn all the layers into training mode
        """
        Module.training(self)
        for m in self.layers:
            m.training()

    def evaluate(self):
        """
        Turn all the layers into evaluate mode
        """
        Module.evaluate(self)
        for m in self.layers:
            m.evaluate()

```

```

In [4]: class Convolutional(Module):
        """
        Convolutional layer
        """
        def __init__(self, inputLength, inputDepth, filterLength, filterDepth):
            Module.__init__(self)
            # Initialization
            stdv = 1./np.sqrt(inputLength)

            self.weight = np.random.uniform(-stdv, stdv, (filterLength,
inputDepth, filterDepth))
            self.gradWeight = np.ndarray((filterLength, inputDepth, filterDepth))
            self.bias = np.random.uniform(-stdv, stdv, filterDepth)
            self.gradBias = np.ndarray(filterDepth)

        def forward(self, _input):
            """
            output = W * filterRegion + b convolution over input

            _input:

```

```

N x inputLength x inputDepth matrix

"""
inputDepth = self.weight.shape[1]
filterDepth = self.weight.shape[2]

_input = _input.reshape(_input.shape[0],-1,inputDepth)

self._output = []
for i in range(filterDepth):
    self._output.append(sg.correlate(_input, np.expand_dims(np.take(self.weight,i,2),0), 'valid') + self.bias[i])
    self._output = np.array(self._output)
    self._output =
self._output.reshape(self._output.shape[1],self._output.shape[2],self._output.shape[0])

    return self._output

def backward(self, _input, _gradOutput):
    """
    _input:
    N x inputLength x inputDepth matrix
    _gradOutputSize:
    N x outputLength x outputDepth matrix
    """

    filterLength = self.weight.shape[0]
    inputDepth = self.weight.shape[1]
    filterDepth = self.weight.shape[2]

    _input = _input.reshape(_input.shape[0],-1,inputDepth)
    _gradOutput = _gradOutput.reshape(_gradOutput.shape[0],-1,filterDepth)

    for i in range(filterDepth):
        self.gradWeight[:, :, i] = sg.correlate(_input, np.expand_dims(np.take(_gradOutput,i,2),2), 'valid')

        self.gradBias = np.sum(_gradOutput, axis=(0,1))

        self._gradInput = []
        for i in range(inputDepth):
            self._gradInput.append(sg.correlate(np.pad(_gradOutput, ((0,0),(filterLength-1,filterLength-1),(0,0)), 'constant'),
            np.expand_dims(np.take(self.weight,i,1),0), 'valid'))
            self._gradInput = np.array(self._gradInput)
            self._gradInput = self._gradInput.reshape(self._gradInput.shape[1],self._gradInput.shape[2],self._gradInput.shape[0])

        return self._gradInput

def parameters(self):
    """
    Return weight and bias and their g
    """

    return [self.weight, self.bias], [self.gradWeight, self.gradBias]

```

```

In [5]: class FullyConnected(Module):
        """
        Fully connected layer
        """
        def __init__(self, inputSize, outputSize):
            Module.__init__(self)
            # Initialization
            stdv = 1./np.sqrt(inputSize)

            self.weight = np.random.uniform(-stdv, stdv, (inputSize, outputSize))
            self.gradWeight = np.ndarray((inputSize, outputSize))
            self.bias = np.random.uniform(-stdv, stdv, outputSize)
            self.gradBias = np.ndarray(outputSize)

        def forward(self, _input):
            """
            output = W * input + b

            _input:
            N x inputSize matrix

            """
            _input = _input.reshape(_input.shape[0],-1)

            self._output = np.dot(_input, self.weight) + self.bias
            return self._output

        def backward(self, _input, _gradOutput):
            """
            _input:
            N x inputSize matrix
            _gradOutputSize:
            N x outputSize matrix
            """
            _input = _input.reshape(_input.shape[0],-1)
            _gradOutput = _gradOutput.reshape(_gradOutput.shape[0],-1)

            self.gradWeight = np.dot(_input.T, _gradOutput)
            self.gradBias = np.sum(_gradOutput, axis=0)

            self._gradInput = np.dot(_gradOutput, self.weight.T)
            return self._gradInput

        def parameters(self):
            """
            Return weight and bias and their g
            """
            return [self.weight, self.bias], [self.gradWeight, self.gradBias]

```



```

In [6]: class ReLU(Module):
        """
        ReLU activation, not trainable.
        """
        def __init__(self):
            Module.__init__(self)
            return

        def forward(self, _input):
            """
            output = max(0, input)

            _input:
            N x d matrix
            """
            #_input = _input.reshape(_input.shape[0], -1)

            self._output = np.maximum(0, _input)
            return self._output

        def backward(self, _input, _gradOutput):
            """
            gradInput = gradOutput * mask
            mask = _input > 0

            _input:
            N x d matrix

            _gradOutput:
            N x d matrix
            """
            #_input = _input.reshape(_input.shape[0], -1)
            #_gradOutput = _gradOutput.reshape(_gradOutput.shape[0], -1)

            self._gradInput = _gradOutput * (_input > 0)
            return self._gradInput

        def parameters(self):
            """
            No trainable parameters, return None
            """
            return None, None

```

```
In [7]: class SoftMaxLoss(object):
        def __init__(self):
            return

        def forward(self, _input, _label):
            """
            Softmax and cross entropy loss layer. Should return a scalar, since i
t's a
            loss. (It's almost identical to what in hw2)

            _input: N x C
            _labels: N x C, one-hot

            Returns: loss (scalar)
            """
            self._output = -np.sum(_label * (_input -
np.log(np.sum(np.exp(_input), axis=1)).reshape(1, -1).T))
            return self._output

        def backward(self, _input, _label):
            self._gradInput = np.exp(_input)/np.sum(np.exp(_input), axis=1).reshap
e(1, -1).T - _label
            return self._gradInput
```

```
In [8]: # Test softmaxloss, the relative error should be small enough
def test_sm():
    crit = SoftMaxLoss()
    gt = np.zeros((3, 10))
    gt[np.arange(3), np.array([1,2,3])] = 1
    x = np.random.random((3,10))
    def test_f(x):
        return crit.forward(x, gt)

    crit.forward(x, gt)

    gradInput = crit.backward(x, gt)
    gradInput_num = numeric_gradient(test_f, x, 1, 1e-6)
    #print(gradInput)
    #print(gradInput_num)
    print(relative_error(gradInput, gradInput_num, 1e-8))

test_sm()
```

4.64713032996e-09

```
In [9]: # Test modules, all the relative errors should be small enough
def test_module(model):

    model.evaluate()

    crit = TestCriterion()
    gt = np.random.random((3,10))
    x = np.random.random((3,10))
    def test_f(x):
        return crit.forward(model.forward(x), gt)

    gradInput = model.backward(x, crit.backward(model.forward(x), gt))
    gradInput_num = numeric_gradient(test_f, x, 1, 1e-6)
    print(relative_error(gradInput, gradInput_num, 1e-8))

# Test fully connected
model = FullyConnected(10, 10)
test_module(model)

# Test ReLU
model = ReLU()
test_module(model)

# Test Sequential
model = Sequential()
model.add(FullyConnected(10, 10))
model.add(ReLU())
#model.add(Dropout())
test_module(model)

4.75853713145e-09
7.36022825478e-10
1.75509770151e-09
```

```

In [10]: # Test gradient descent, the loss should be lower and lower
trainX = np.random.random((10,25,15))

model = Sequential()
model.add(Convolutional(25,15, 6,20))
model.add(ReLU())
model.add(Convolutional(20,20, 11,5))
model.add(FullyConnected(50,1))

crit = TestCriterion()

it = 0
state = None
while True:
    output = model.forward(trainX)
    loss = crit.forward(output, None)
    if it % 100 == 0:
        print(loss)
    doutput = crit.backward(output, None)
    model.backward(trainX, doutput)
    params, gradParams = model.parameters()
    sgdmmom(params, gradParams, 0.0005, 0.8)
    if it > 1000:
        break
    it += 1

0.411563871173
0.15327683183
0.117801346354
0.0804396823411
0.0675119314006
0.051551482172
0.0353256828849
0.0225411762008
0.0234306934355
0.0212756206
0.0122138643986

```

Now we start to work on real data.

```
In [11]: import MNIST_utils
data_fn = "CLEAN_MNIST_SUBSETS.h5"

# We only consider large set this time
print("Load large trainset.")
Xlarge,Ylarge = MNIST_utils.load_data(data_fn, "large_train")
print(Xlarge.shape)
print(Ylarge.shape)

print("Load valset.")
Xval,Yval = MNIST_utils.load_data(data_fn, "val")
print(Xval.shape)
print(Yval.shape)
```

```
Load large trainset.
(7000L, 576L)
(7000L, 10L)
Load valset.
(2000L, 576L)
(2000L, 10L)
```

```

In [12]: def predict(X, model):
          """
          Evaluate the soft predictions of the model.
          Input:
          X : N x d array (no unit terms)
          model : a multi-layer perceptron
          Output:
          yhat : N x C array
                  yhat[n][:] contains the score over C classes for X[n][:]
          """
          return model.forward(X)

def error_rate(X, Y, model):
    """
    Compute error rate (between 0 and 1) for the model
    """
    model.evaluate()
    res = 1 - (model.forward(X).argmax(-1) == Y.argmax(-1)).mean()
    model.training()
    return res

from copy import deepcopy

def runTrainVal(X,Y,model,Xval,Yval,trainopt):
    """
    Run the train + evaluation on a given train/val partition
    trainopt: various (hyper)parameters of the training procedure
    During training, choose the model with the lowest validation error. (early
    stopping)
    """

    eta = trainopt['eta']

    N = X.shape[0] # number of data points in X

```

```

# Save the model with lowest validation error
minValError = np.inf
saved_model = None

shuffled_idx = np.random.permutation(N)
start_idx = 0
for iteration in range(trainopt['maxiter']):
    if iteration % int(trainopt['eta_frac'] * trainopt['maxiter']) == 0:
        eta *= trainopt['etadrop']
        # form the next mini-batch
        stop_idx = min(start_idx + trainopt['batch_size'], N)
        batch_idx = range(N)[int(start_idx):int(stop_idx)]
        bX = X[shuffled_idx[batch_idx],:]
        bY = Y[shuffled_idx[batch_idx],:]

        score = model.forward(bX)
        loss = crit.forward(score, bY)
        # print(loss)
        dscore = crit.backward(score, bY)
        model.backward(bX, dscore)

        # Update the data using
        params, gradParams = model.parameters()
        sgd(params, gradParams, eta, weight_decay = trainopt['lambda'])
        start_idx = stop_idx % N

    if (iteration % trainopt['display_iter']) == 0:
        #compute train and val error; multiply by 100 for readability (make it percentage points)
        trainError = 100 * error_rate(X, Y, model)
        valError = 100 * error_rate(Xval, Yval, model)
        print('{:8} batch loss: {:.3f} train error: {:.3f} val error: {:.3f}'.format(iteration, loss, trainError, valError))

        if valError < minValError:
            saved_model = deepcopy(model)
            minValError = valError

return saved_model, minValError, trainError

```

```

In [13]: def build_model(input_size, hidden_size, output_size, activation_func =
          'ReLU', dropout = 0):
          """
          Build the model:
          input_size: the dimension of input data
          hidden_size: the dimension of hidden vector
          output_size: the output size of final layer.
          activation_func: ReLU, Logistic, Tanh, etc. (Need to be implemented by you
rself)
          dropout: the dropout rate: if dropout == 0, this is equivalent to no dropo
ut
          """
          model = Sequential()
          model.add(Convolutional(input_size,1, hidden_size,1))
          model.add(FullyConnected((input_size-hidden_size+1), output_size))
          model.add(ReLU())
          model.add(FullyConnected(output_size, output_size))

          return model

```



```

In [14]: # -- training options
trainopt = {
    'eta': .001,    # initial learning rate
    'maxiter': 2500, # max number of iterations (updates) of SGD
    'display_iter': 500, # display batch loss every display_iter updates
    'batch_size': 100,
    'etadrop': .5, # when dropping eta, multiply it by this number (e.g., .5 means halve it)
    'eta_frac': .25 #
}

NFEATURES = Xlarge.shape[1]

# we will maintain a record of models trained for different values of lambda
# these will be indexed directly by lambda value itself
trained_models = dict()

# set the (initial?) set of lambda values to explore
lambdas = np.array([0, 0.001, 0.01, 0.1])
hidden_sizes = np.array([10])

for lambda_ in lambdas:
    for hidden_size_ in hidden_sizes:
        trainopt['lambda'] = lambda_
        model = build_model(NFEATURES, hidden_size_, 10, dropout = 0)
        crit = SoftMaxLoss()
        # -- model trained on large train set
        trained_model, valErr, trainErr = runTrainVal(Xlarge, Ylarge, model, Xval, Yval, trainopt)
        trained_models[(lambda_, hidden_size_)] = {'model': trained_model, "val_err": valErr, "train_err": trainErr }
        print('train set model: -> lambda= %.4f, train error: %.2f, val error: %.2f' % (lambda_, trainErr, valErr))

best_trained_lambda = 0.
best_trained_model = None
best_trained_val_err = 100.
for (lambda_, hidden_size_), results in trained_models.items():
    print('lambda= %.4f, hidden size: %5d, val error: %.2f' % (lambda_, hidden_size_, results['val_err']))
    if results['val_err'] < best_trained_val_err:
        best_trained_val_err = results['val_err']
        best_trained_model = results['model']
        best_trained_lambda = lambda_

print("Best train model val err:", best_trained_val_err)
print("Best train model lambda:", best_trained_lambda)

```